

# ECE 598NSG/498NSU

## Deep Learning in Hardware

### Fall 2020

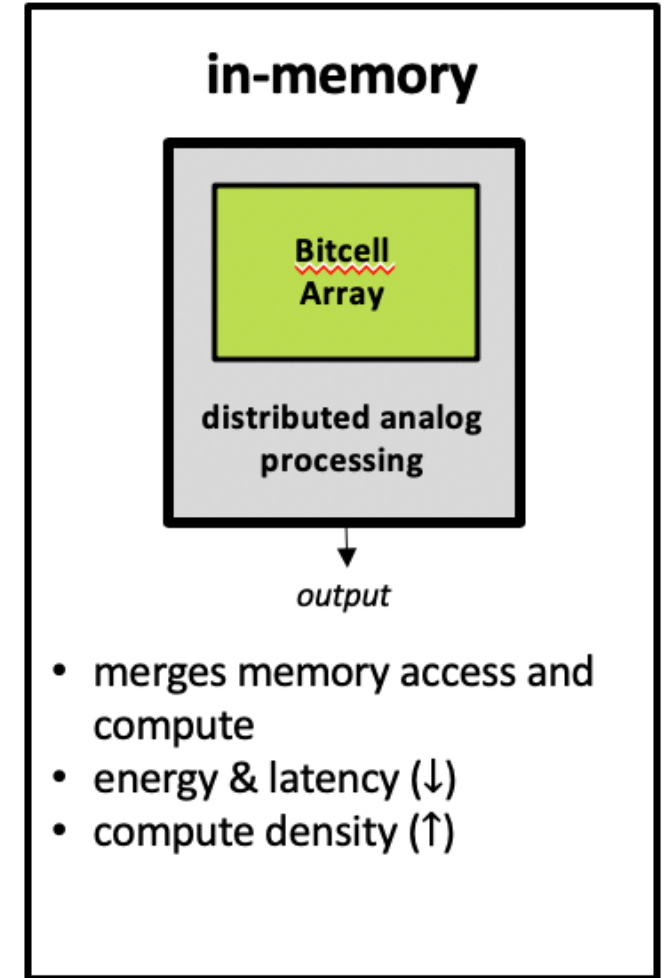
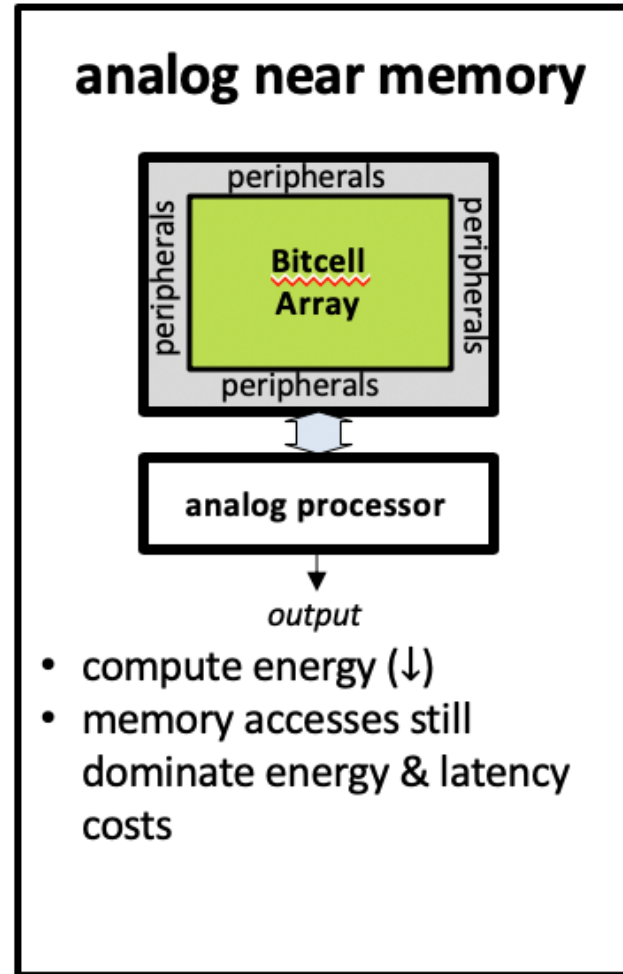
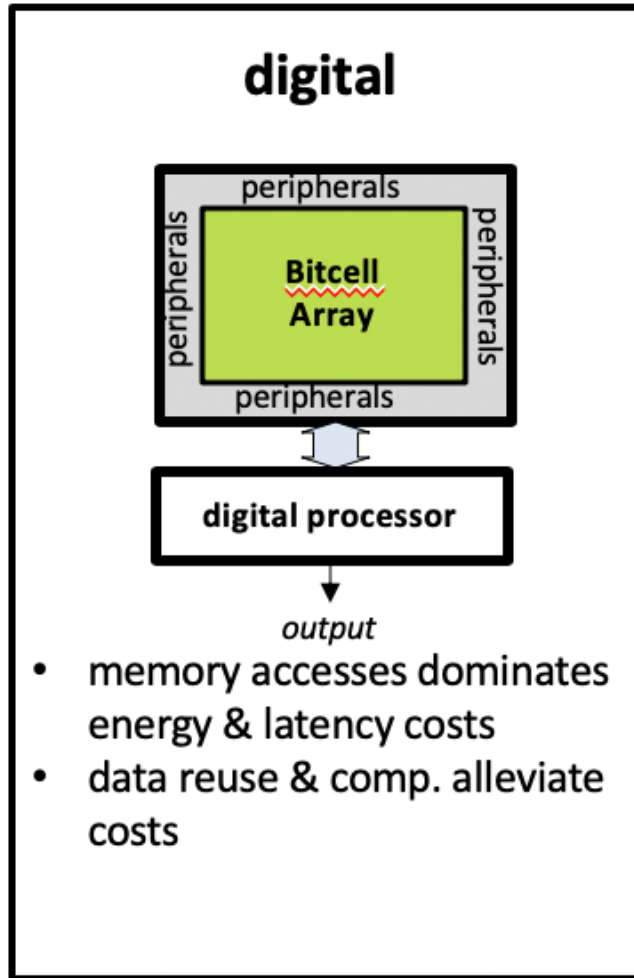
DNN Accelerators – Overview  
Roofline, Floorline plots, Reuse

Naresh Shanbhag

Department of Electrical and Computer Engineering  
University of Illinois at Urbana-Champaign


<http://shanbhag.ece.uiuc.edu>

# Architectures for Inference




# DNN Accelerators

TPU V1




inference only  
8b fixed-point

TPU V2



inference & training  
fixed-point & floating-point



TPU V3



inference & training  
fixed-point & floating-point

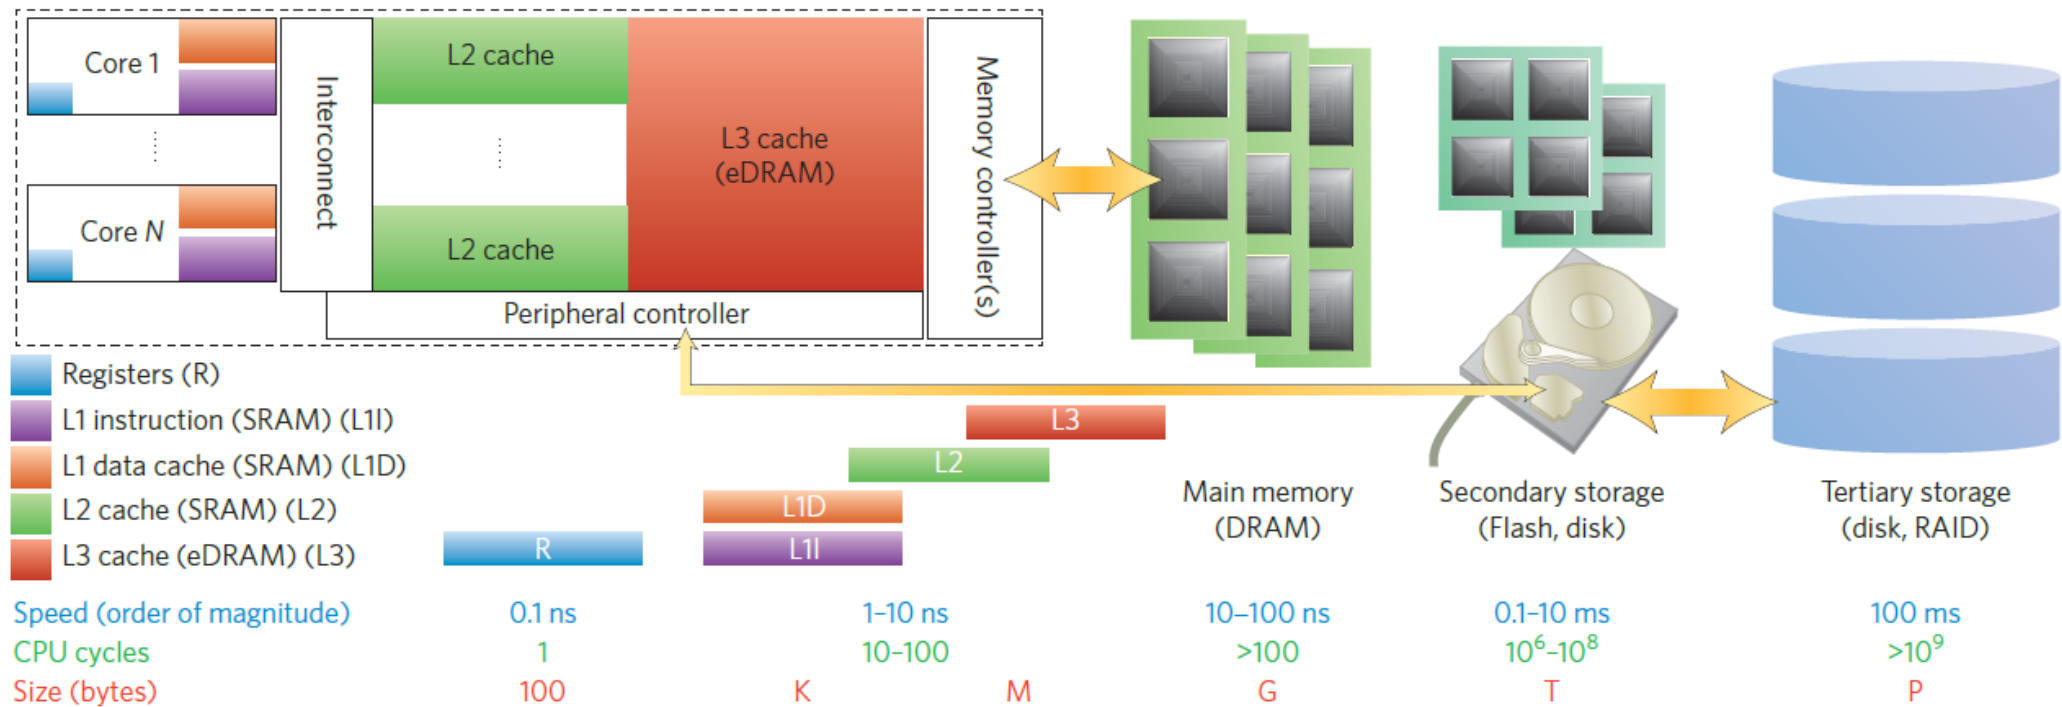
- going beyond GPUs
- accelerating inference & training in the cloud
- intersection of workload, architecture, circuits, software

# Challenge: Memory Access Energy in DNNs



$$E_R = \frac{\text{energy/access}}{\text{energy/OP}}$$


- memory access-to-compute energy consumption ratio
- SRAM accesses:  $E_R = 100\times$  ; DRAM accesses:  $E_R = 1000\times$



[Wong, Salahuddin, Nature Nanotechnology, 2015]


- deep memory hierarchy
- high density → more latency, energy
- **data movement** → dominant cause of energy dissipation in machine learning applications

the memory funnel



# Power Breakdown

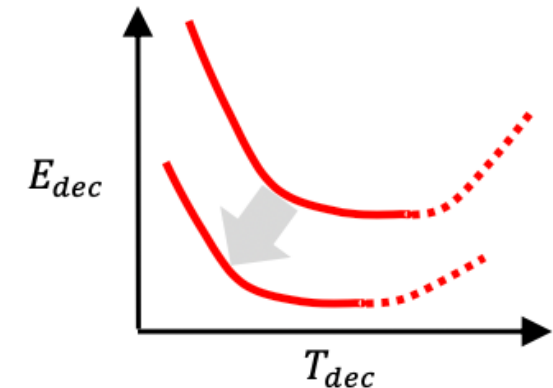
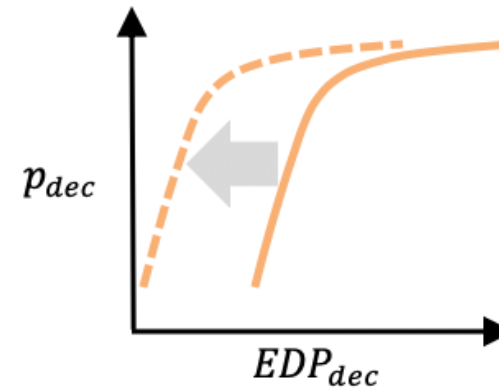

[Eyeriss, Y.-H. Chen, JSSC 2017]



- Data movement costs 45% of total power
- Computation costs about 10%
- Clock network = 33% (hidden cost of digital realizations)

# System Metrics for Inference

system-level **energy vs. latency vs. accuracy** trade-off



- ❑  $T_{dec}$ : decision latency
- ❑  $p_{dec}$ : decision accuracy
- ❑  $EDP_{dec} = E_{dec} \times T_{dec}$ : decision EDP
- ❑  $f_{dec}$ : decision throughput

- ❑  $E_{dec} \propto 1/T_{dec} \propto \text{data \& network size}$
- ❑  $p_{dec} \propto EDP_{dec}$


# Outline

- DNN architecture abstraction
- visualization via the roofline & floorline plots
- data reuse



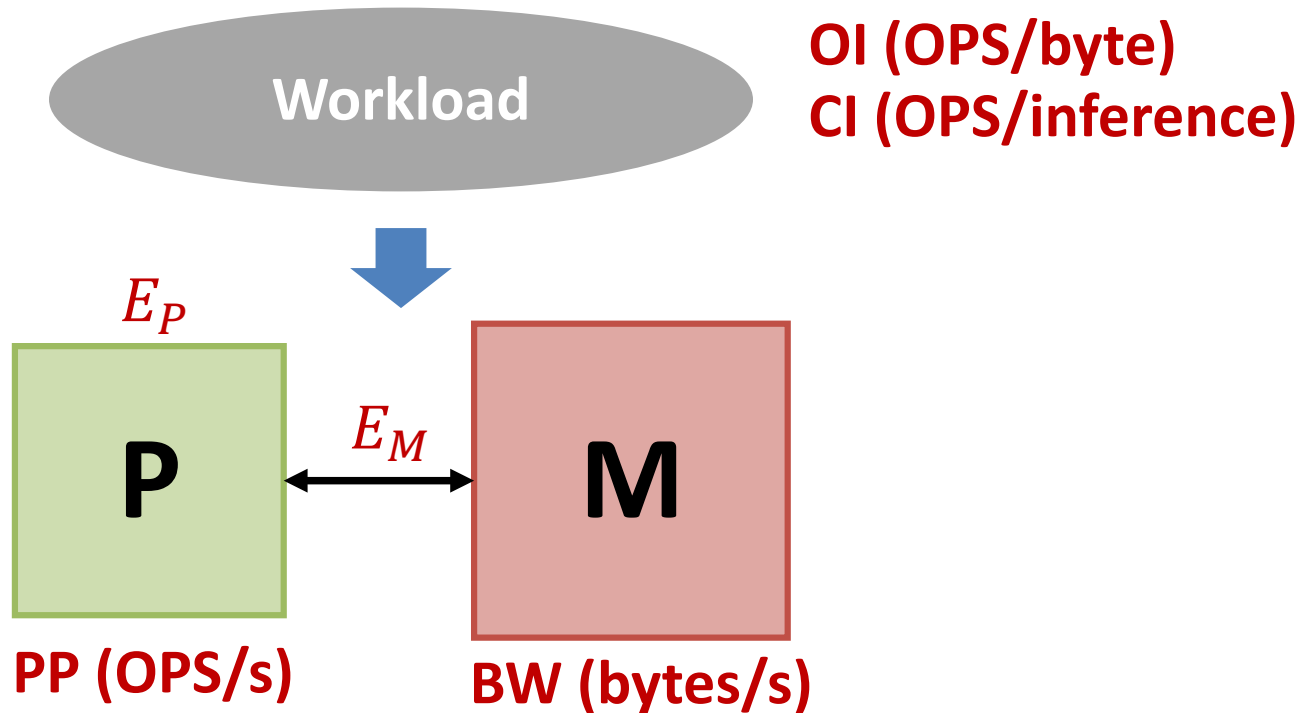
# DNN Architectures – An Abstract View

# DNN Architecture - Abstract View




- Workload – DNNs, CNNs, LSTMs,....
- Processor **P** – GPU, accelerators, many-core architectures,....
- Memory **M** – register, SRAM, DRAM, flash,.....

# Workload Parameters




- operational intensity (OI): # of operations/byte (high is better) – reuse factor
- computational intensity (CI): # of operations per decision (small is better)

# Processor Parameters




- peak performance (PP): # of operations/s (e.g., TOPS =  $10^{12}$  OPS/s)
- arithmetic efficiency ( $E_P$ ): energy per operation (e.g.  $E_P = 0.2pJ/op$ )

# Memory Parameters




- read bandwidth (BW): # of bytes read/s (higher the better)
- energy/read ( $E_M$ ): energy per byte (lower the better)
- $E_R = \frac{E_M}{E_P} = 100$  (SRAM);  $= 500$  (DRAM);  $= 1000$  (flash)

# System Level Metrics




- useful for system-level benchmarking
- energy/decision ( $E_{dec}$ )
- decision throughput ( $f_{dec} = \frac{1}{T_{dec}}$ )

# Component Level Metrics



- used for hardware benchmarking
- energy/OP ( $E_P$ ); peak OP/s ( $PP$ ); bytes read/s ( $BW$ ); energy/byte ( $E_M$ )

# System vs. Component Level Metrics


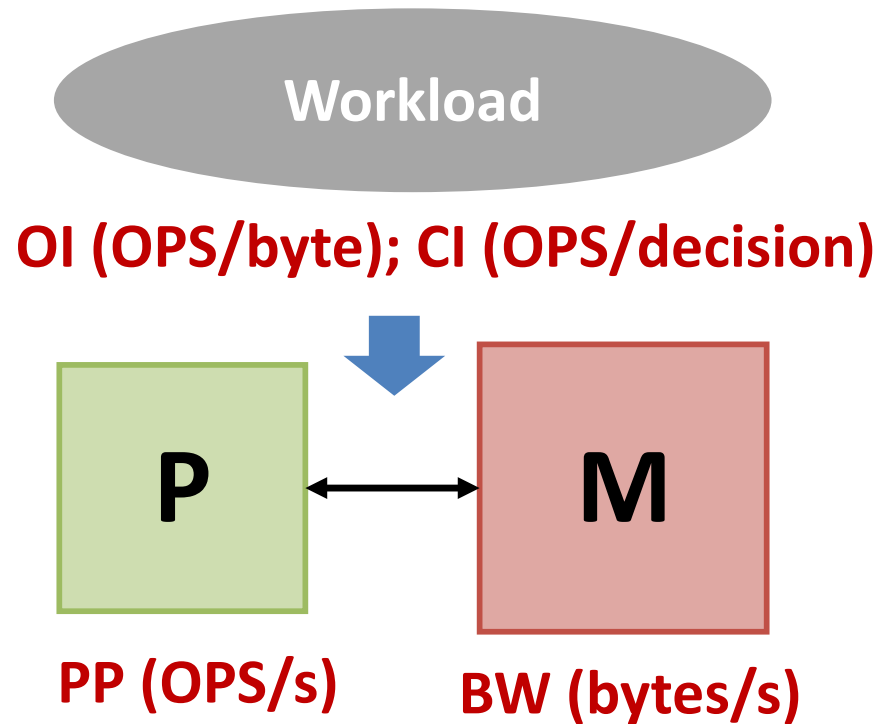


- system-level metrics are a function of component-level metrics, workload requirements, and workload mapping onto hardware
- system-level metrics are the true metrics – hard to measure and report
- alternative – measure @ the component-level and estimate @ the system-level



# Roofline & Floorline Plots – Visual Performance & Efficiency Models

# Architectural-level Performance - Roofline Plot




$$AP = \min(PP, OI \times BW)$$

- $AP$ : attainable performance of the architecture for the workload(s)

- Note:  $AP \leq PP$

(Williams, Waterman, Patterson, Comm. of ACM '09)


# System-level Performance - Bytes/Decision - $M_{dec}$




- Memory bandwidth (BW)  $\neq$  bytes/decision ( $M_{dec}$ )
- Given a workload  $OI$  and  $CI$ , bytes/decision :

$$M_{dec} = \frac{\text{bytes}}{\text{decision}} = \frac{\text{number of OPs/dec}}{\text{number of OPs/byte}} = \frac{CI}{OI}$$

Roofline plot




# System-level Performance - Decision Throughput ( $DT$ )



- Attainable Performance (AP)  $\neq$  Decision throughput (DT)
- Given a workload  $OI$  and  $CI$ , decision throughput (dec/s):

$$DT = \frac{\text{decisions}}{s} = \frac{\text{achievable OPs/s}}{\# \text{ of OPs/dec}} = \frac{AP}{CI} = \min \left[ \frac{PP}{CI}, OI \times \left( \frac{BW}{CI} \right) \right] = \min \left[ \frac{PP}{CI}, \frac{BW}{M_{dec}} \right]$$




# System Energy Efficiency – Floorline Plot

(Shanbhag, Dbouk, Sakr, ECE 498NSU/ECE 598NSG 2019)

- Decision energy (pJ/dec):


$$E_{dec} = CI \times E_P + M_{dec} \times E_M = CI \times E_P + \frac{CI}{OI} \times E_M = CI \left( E_P + \frac{E_M}{OI} \right) = CI \times E_P \left( 1 + \frac{E_R}{OI} \right)$$

Floorline plot



$$E_R = \frac{\text{energy/access}}{\text{energy/OP}} \sim 10^2 \text{ (SRAM)}$$

# Maximizing Inference Throughput

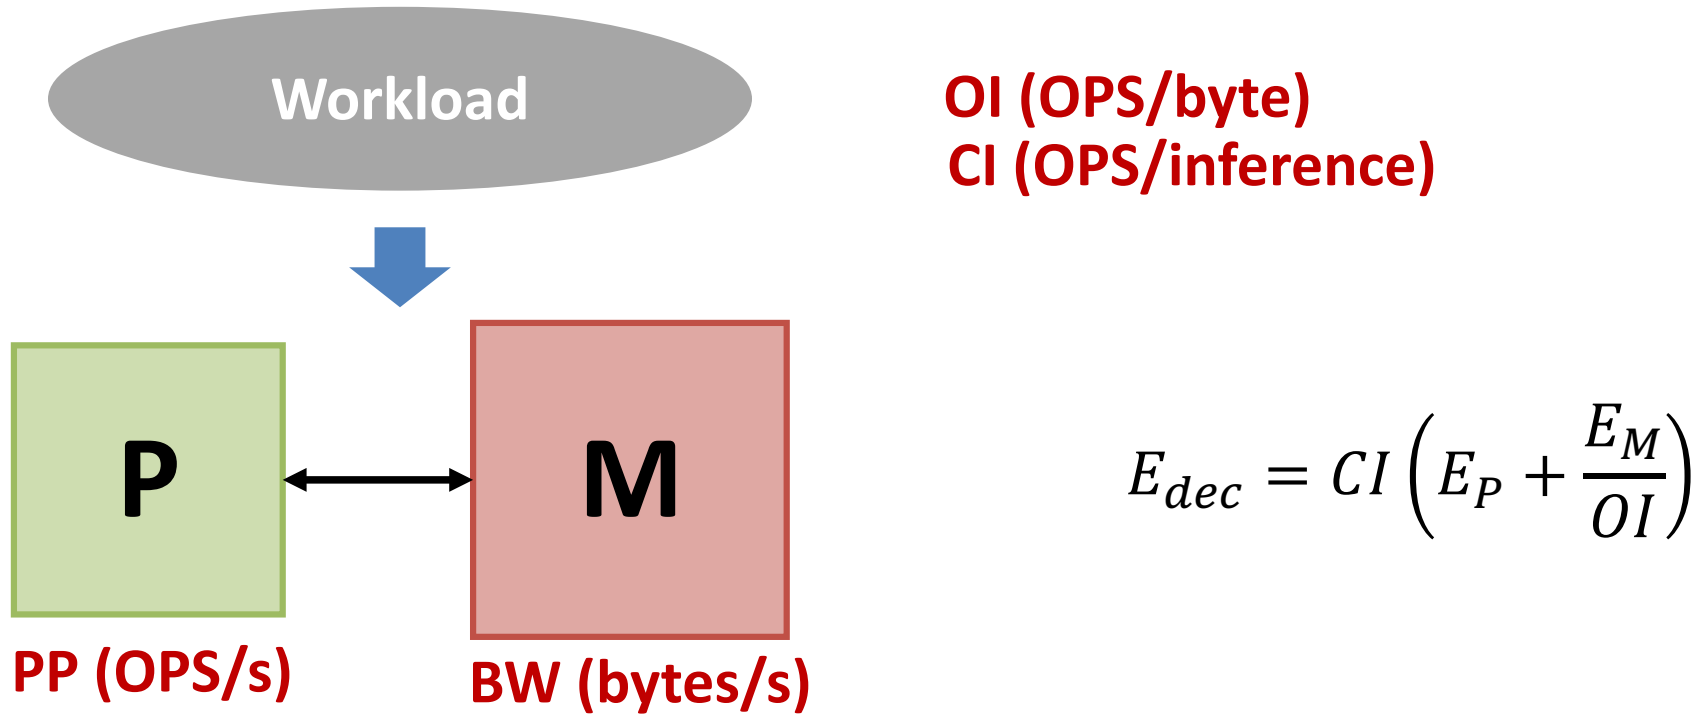


OI (OPS/byte)  
CI (OPS/inference)

$$DT = \frac{AP}{CI} = \frac{\min(PP, OI \times BW)}{CI}$$

- Increase OI: data reuse & HW mapping → helps reduce decision energy
- Decrease CI: low-complexity DNNs, fast algorithms (FFT) → helps reduce decision energy
- Increase PP: systolic architectures & algorithm transforms
- Increase BW: new memory technologies + in-memory computing

# Maximizing Energy Efficiency



- Increase OI: data reuse & HW mapping → helps increase decision throughput
- Decrease CI: low-complexity DNNs → helps increase decision throughput
- Decrease  $E_P$ : energy-efficient computing (e.g. near- $V_t$  computing), reduce precision
- Decrease  $E_M$ : new memory technologies + in-memory computing

# Example


Given:

- an inference workload:
  - $CI = 1000$  OPs/dec
  - $OI = 10$  OPs/byte
- hardware platform with:
  - $PP = 256$  OPs/s
  - $BW = 16$  bytes/s
  - $E_P = 0.23$  pJ/OP
  - $E_M = 2.5$  pJ/byte

Application



Hardware



determine decision energy and latency



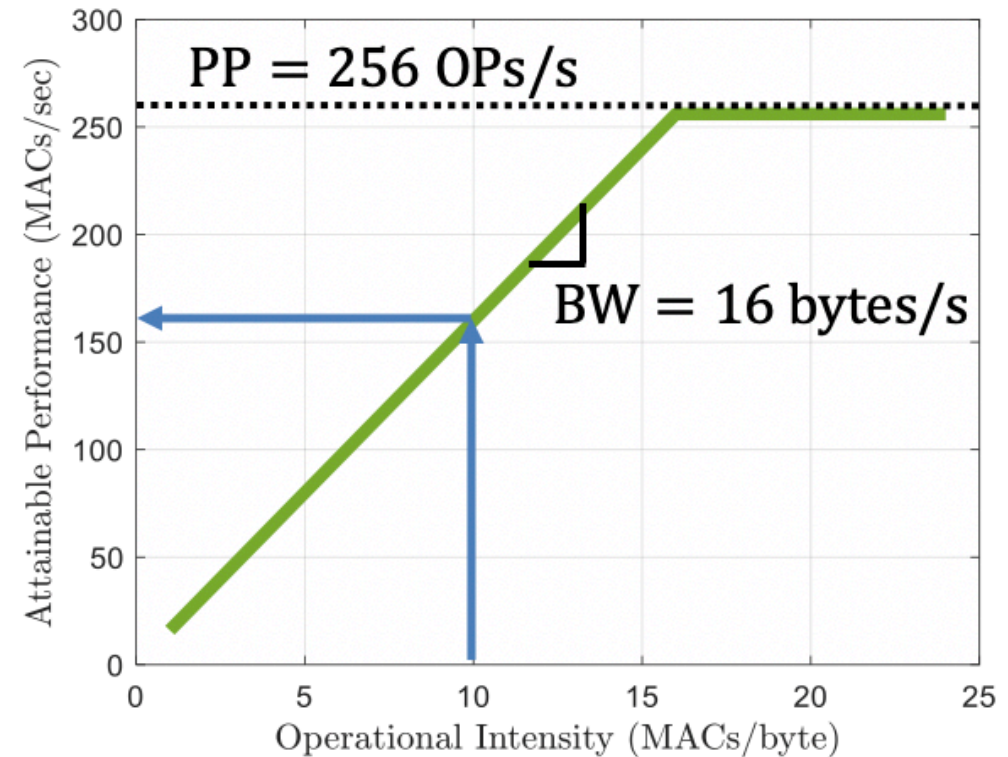
# Throughput - Roofline plot

- Attainable performance:

$$AP = \min(PP, OI \times BW) = 160 \text{ OPs/s}$$

- Decision throughput:

$$DT = \frac{AP}{CI} = \frac{160}{1000} = 0.16 \text{ dec/s}$$



# Energy Efficiency - Floorline plot

- Decision energy:


$$CI \left( E_P + \frac{E_M}{OI} \right) = 1000 \left( 0.23 + \frac{2.5}{10} \right) = 0.48 \text{ nJ/dec}$$

- Floorline plot:

(lower bound)  $CI \times E_P = 0.23 \text{ nJ}$

$$E_R = \frac{E_M}{E_P} = 10.87$$

$$E_{\text{dec}}(\text{nJ}) = 0.23 \left( 1 + \frac{10.87}{OI} \right)$$




# Component-level Energy-Delay Trade-offs

- Energy/OP ( $E_P$ ) & Peak OP/s ( $PP$ ) – low-power and high-throughput digital architectures and circuits (high-ceiling home)
- Bytes read/s ( $BW$ ) & Energy/byte ( $E_M$ ) – energy-efficient memory and interface design
- some of these are mature topics but keep evolving to meet the needs of machine learning applications
- energy consumption vs. delay (vs. robustness) trade-off
- energy vs. delay trade-off captured via **energy-delay product**  $EDP$
- energy/delay vs. robustness trade-off is more subtle (future lecture)

# Roofline Plot – A Visual Performance Model

# Example – Memory Bound Scenario

[adapted from Hwu – IBM AI Compute Symposium, 2019]




SP: single-precision (32-bit)

Each operand must be used **62.3 times** (**reuse factor**) once fetched to achieve peak FLOPS rate (**ridge point**)

or

Sustain < **1.6%** of peak without data reuse

$$\text{Achievable Performance (AP)} = \min(\text{PP}, \text{OI} \times \text{BW})$$



# Roofline: An Insightful Visual Performance Model for Floating-Point Programs and Multicore Architectures\*


Samuel Williams, Andrew Waterman, and David Patterson

Parallel Computing Laboratory, 565 Soda Hall, U.C. Berkeley, Berkeley, CA 94720-1776, 510-642-6587  
samw, waterman, pattnsn@eecs.berkeley.edu

*“A model need not be perfect, just insightful.”*

- Motivation: off-chip memory bandwidth will limit the performance
- Need a simple model relates processor performance to off-chip memory traffic similar to Amdhal's Law
- **Amdhal's Law:** performance gain of a parallel computer is limited by the serial portion of the workload

# Roofline for AMD Opteron X2




[Williams, ACM Communications, 2019]

- 2 core processor
- BW = 15 GB/s; PP = 17.6 GFlops/s




# Opteron X2 vs. X4



- X2 = 2 core (2.2 GHz) X4 = 4 core (2.3 GHz)
- workloads with  $> 1$  Flop/Byte OI needed for X4


# Roofline Plots – The Reality

- Roofline plot provides **a ceiling** on the achievable performance
- They ignore **intricate data movement** required by applications
- Different layer shapes might cause idle PEs => unutilized resources



[TPU – ISCA'17]

# Dot Product Example



- Processor **P** with a peak performance (PP) of 256 operations/sec
- Memory **M** with a bandwidth (BW) of 16 bytes/sec
- Workload: perform a 1D “convolution” (correlator):

$$y_i = \sum_{j=1}^M w_j x_{i+j-1}$$

where  $x \in \mathbb{R}^N$ ,  $w \in \mathbb{R}^M \Rightarrow y \in \mathbb{R}^{N-M+1}$  (with zero-padding), and  $M \in [N]$   
assume 32b floating point operations;  $N = 128$ ;

# Maximum AP

- Assume we only have to fetch  $w$  from  $\mathbf{M}$ .
- OI for 1D convolution (DP):


$$OI = \frac{\text{number of MACs}}{\text{number of bytes read}} = \frac{\overbrace{(N - M + 1) \times M}^{\text{number of dot products}}}{4 \times M} \rightarrow \text{\# MACs/DP}$$

$$OI = \frac{(N - M + 1)}{4}$$


- OI **reduces** with kernel size  $M$
- OI **increases** with input data size  $N$

each weight value (32b)  
needs 4 memory reads

# Varying Kernel Length ( $M$ )




# Doubling Memory $BW$




- $M = 32$ :  $AP$  is *compute bound*  $\rightarrow$  no benefit from doubling  $BW$
- $M = 96$ ;  $AP$  is *memory bound*  $\rightarrow$  benefits from doubling  $BW$

# Doubling Peak Performance ( $PP$ )



- $M = 32$ :  $AP$  increases and it becomes *memory bound*
- $M = 96$ :  $AP$  is already *memory bound*, so no change


# Methods to Improve AP



- A → B: same architecture - **enhance reuse**
- A → C: **same reuse** & processor – enhance BW
- B → D: same interface – **enhance reuse** and PP
- A → E: tunable - **enhance reuse**, BW and PP



# Data Reuse



OI (OPS/byte)  
CI (OPS/inference)

$$DT = \frac{AP}{CI} = \min \left[ \frac{PP}{CI}, OI \times \left( \frac{BW}{CI} \right) \right]$$

$$E_{dec} = CI \left( E_P + \frac{E_M}{OI} \right)$$

- **Increase OI: data reuse** & HW mapping
- Decrease CI: low-complexity DNNs
- Increase PP: systolic architectures & algorithm transforms
- Increase BW: new memory technologies + in-memory computing


# Terminology

- **Data reuse**: # of MACs per data value **read** (i.e., MACs/data).
- For every MAC operation involved in a convolution operation, there are 3 data types:

$$\text{output partial sum} \longleftarrow z \leftarrow z + \overset{\text{weight}}{\uparrow} w \times x \longrightarrow \text{input activation}$$

- Hence, three types of data reuse: input, weight, and partial sum
- Partial sum reuse = DP length = # of partial sums needed to compute one DP

# 2D Convolution




# Input Reuse



- Each **input pixel** is **reused  $\sim K \times K$  times** across the same weight filter
- Each **input pixel** is **reused  $M$  times** due to  $M$  filters processing the same input feature
- Total input reuse:  $K^2 M$
- e.g., Layer 1 of VGG-16 on ImageNet –  $3 \times 3 \times 64 = 576$  MACs/byte = OI

# Weight Reuse



- Each **weight pixel** is **reused  $E \times F$  times** across the same input feature map
- Each **weight pixel** is **being used  $N$  times** when multiple ( $N > 1$ ) input batches are processed
- Total weight reuse:  $E \times F \times N$
- e.g., Layer 1 of VGG-16 on ImageNet –  $224 \times 224 \times 1 = 50176$  MACs/byte = OI

## Course Web Page

<https://courses.grainger.illinois.edu/ece598nsg/fa2020/>

<https://courses.grainger.illinois.edu/ece498nsu/fa2020/>

<http://shanbhag.ece.uiuc.edu>